



Kognit Validated

TECHNICAL DNA

Python

C++

JavaScript

HTML

FRAMEWORKS

Textual

SDL2

OpenGL

Flask

Express

ONNX Runtime

FOCUS

Architecture Astronautics and Buzzword-Driven Development

INFERRED PERSONA

Architecture Astronaut and Professional Over-Engineer

SOURCES

<https://github.com/Pomilon>

<http://pomilon.xyz>

Pomilon

Architecture astronaut building a graveyard of over-engineered AI toys with zero users

Executive Summary

Pomilon is the quintessential architecture astronaut - a developer who mistakes complexity for sophistication and ships nothing but elaborate technical debt. With a bio that reads like a LinkedIn humblebrag ("Tech nerd and hobbyist, I try building whatever idea comes to my mind"), this developer has managed to accumulate a portfolio of 12 repositories, each more tragically over-engineered than the last.

The technical schizophrenia is immediately apparent: oscillating between C++ systems programming and Python AI wrappers like a caffeinated intern who just discovered both Docker and language models. Their pinned projects read like a buzzword bingo card - "Continuous Reasoning State Model," "Mamba-MoE Language Model," "Hierarchical State Sovereignty" - each phrase more meaningless than the last.

What's particularly galling is the commitment to building everything from scratch. While normal developers would use existing package managers, Pomilon built "Peck" for their custom language "Pome" - because apparently pip, npm, cargo, and the other 47 package managers weren't quite pretentious enough. The 2 followers (probably their mom and their alt account) are undoubtedly impressed by the Docker-compose files that could orchestrate a small data center for projects that, judging by the commit history, have never been used by a single human being.

The crown jewel of this technical theater is Plexir – a "modular, keyboard-centric AI terminal workspace" that combines "multi-provider LLM orchestration" with "persistent Docker sandboxing." Translation: it's a over-engineered CLI wrapper that can call OpenAI's API while running in Docker, because apparently `curl` was too simple. The complexity score of 8/10 for what amounts to a fancy chat client is either delusional brilliance or Stockholm syndrome.

This is a developer who has mastered the art of building solutions to problems that don't exist, using technologies that won't exist by the time anyone might actually need them.

Key Projects & Impact

Plexir

Developer • Python, Textual, Docker, MCP, OpenAI API, Gemini API, Groq API

An over-engineered TUI that orchestrates multiple LLM APIs because apparently using one API at a time wasn't complicated enough

Successfully achieved a complexity score of 8/10 for what's essentially a fancy curl wrapper

CRSM

Developer • Python, Mamba, State Machines

An asynchronous 'System 2' architecture implementing 'Hierarchical State Sovereignty' in a Mamba backbone

Managed to make state machines sound like a geopolitical crisis

Peck

Developer • C++, Package Management, Dependency Resolution

A package manager for a language nobody uses, solving problems solved decades ago

Achieved 1 star, presumably from the developer's alt account

Technical Deep Dive

Technical Deep Dive: The Architecture Astronaut's Masterpiece

The Pomilon Paradox: Building Skyscrapers in a Desert

Pomilon represents a fascinating case study in what happens when technical capability completely divorces from practical application. This developer has achieved something remarkable: creating some of the most technically complex solutions to non-existent problems I've ever encountered. Let's dissect the architectural carnage.

Plexir: The \$50M Startup That Nobody Asked For

Plexir is a masterclass in over-engineering. At its core, it's a TUI (Text User Interface) wrapper for LLM APIs, but the developer managed to turn this simple concept into a distributed systems nightmare. The architecture includes:

Multi-Provider Orchestration Layer: Instead of just calling OpenAI's API directly, Plexir implements a sophisticated provider abstraction that can route requests to OpenAI, Gemini, Groq, or any other LLM provider. This includes automatic fallback mechanisms, rate limiting, and quota management – because apparently API reliability is such a massive problem that we need circuit breakers for text generation.

Persistent Docker Sandboxing: Every LLM interaction runs inside a Docker container that persists between sessions. The justification is "security and isolation," but in reality, it's using a sledgehammer to crack a nut. The containers are managed through a custom orchestration layer that handles container lifecycle, resource limits, and cleanup – essentially recreating a subset of Kubernetes for running Python scripts.

MCP Protocol Implementation: The Model Context Protocol integration is particularly galling. MCP is designed to standardize how applications interact with LLMs, but Plexir implements it in a way that adds layers of abstraction without providing any actual standardization. The protocol layer includes custom serialization formats, authentication mechanisms, and error handling – all for what's essentially JSON over HTTP.

Plugin Architecture: The modular plugin system allows extending functionality through dynamically loaded modules. Each plugin runs in its own isolated context with dependency injection, configuration management, and lifecycle hooks. The complexity here is staggering – we're talking about a plugin system for a TUI application that most people would implement as a simple Python module.

Context Management System: Plexir maintains conversation history across sessions using SQLite with JSON schema validation. The context system includes conversation branching, merging, and pruning algorithms that would be impressive if they served any purpose beyond making simple state management unnecessarily complex.

The real tragedy is that all of this complexity exists to solve problems that don't exist. API reliability? Most LLM providers have 99.9% uptime. Security? The containers are running locally with the same permissions as the host user. Context management? Most users just want to chat with an AI, not manage a distributed conversation graph.

Kestr: Reinventing grep with Machine Learning

Kestr is positioned as a "high-performance daemon for real-time codebase indexing," but it's actually an elaborate exercise in using machine learning to solve problems that grep solved decades ago.

Real-time File System Monitoring: The daemon uses inotify to watch for file system changes, triggering immediate re-indexing. The implementation includes custom debouncing algorithms, change batching, and priority queues to handle high-frequency updates. For most codebases, this is massive overkill – files don't change that frequently, and when they do, a simple delay-and-batch approach would work fine.

Semantic Embedding Generation: Instead of simple text indexing, Kestr generates semantic embeddings for code using ONNX Runtime. The embedding model is a custom transformer trained on code that converts source code into 768-dimensional vectors. The rationale is to enable "semantic search" – finding code based on meaning rather than text matching. In practice, this means you can't find a function by name unless the embedding model thinks it's semantically similar to your query.

HNSW Vector Search: The embeddings are stored in an HNSW graph for similarity search. The implementation includes custom graph construction algorithms, search heuristics, and memory management. The graph supports incremental updates,

distributed queries, and approximate search with tunable recall/latency tradeoffs. For a codebase that probably has fewer than 10,000 files, this is like using a Formula 1 car to deliver pizza.

MCP Integration: Like Plexir, Kestr implements MCP for "AI agent compatibility." The integration includes custom protocol handlers, authentication, and rate limiting. The idea is that AI agents can query the codebase index using natural language, but in reality, it just adds another layer of complexity to what could be a simple REST API.

Performance Optimizations: The codebase includes memory-mapped file I/O, parallel processing with `std::async`, custom memory allocators, and CPU cache-aware data structures. These optimizations might make sense for indexing terabytes of data, but for a typical project with a few thousand files, they're just premature optimization theater.

The architectural decisions here reveal a developer who's read about vector databases and decided that every problem needs machine learning. The reality is that most code search problems are solved perfectly well by `grep`, `ripgrep`, or simple text indexing. Adding semantic embeddings just makes the system less reliable and harder to debug.

CRSM: Making State Machines Complicated Again

CRSM (Continuous Reasoning State Model) is described as an "asynchronous 'System 2' architecture that implements Hierarchical State Sovereignty within a Mamba backbone." Let's translate this buzzword salad:

Mamba Backbone: The system uses Mamba (a state space model) as the core computational engine. Mamba is designed for sequence modeling, but here it's being used to implement state transitions in a finite state machine. It's like using a neural network to implement if statements.

Hierarchical State Sovereignty: This appears to be a fancy way of saying "nested state machines." The hierarchy allows states to contain sub-states, with complex transition rules between levels. The "sovereignty" concept means that each state has complete control over its internal behavior and can reject transitions from parent states. In practice, this creates a Byzantine system where state transitions can fail for arbitrary reasons.

Asynchronous System 2 Architecture: The "System 2" reference is to dual-process theory in psychology – System 1 is fast and intuitive, System 2 is slow and deliberative. CRSM implements "System 2" reasoning by making all state transitions asynchronous and adding deliberation steps. This means that changing state requires multiple asynchronous operations, making the system slow and unpredictable.

Continuous Reasoning: The system continuously re-evaluates state transitions even when no events occur. This creates a background process that's constantly running Mamba inference to check if the current state is still valid. For a state machine, this is completely unnecessary – states should only change in response to events, not continuous re-evaluation.

The implementation includes custom async runtime, state persistence, distributed consensus, and failure recovery. For a state machine. A simple switch statement would have solved the same problem with 1/100th the complexity.

Peck: The Package Manager That Time Forgot

Peck is a package manager for Pome (a custom scripting language), recreating 30 years of package management history with the added complexity of being written in C++.

Dependency Resolution Algorithm: Peck implements a constraint satisfaction solver for dependency resolution that handles diamond dependencies, version conflicts, and circular dependencies. The algorithm uses backtracking with heuristics, conflict-driven learning, and parallel search. For a language with zero packages, this is like building a nuclear reactor to power a light bulb.

Virtual Environment System: The virtual environment implementation creates isolated package installations with custom filesystem virtualization, environment variable management, and process isolation. The system includes a custom shell integration that activates environments automatically based on directory. For a language that nobody uses, this is isolation theater.

Package Index System: Peck maintains a separate package index repository that requires its own update mechanism, mirroring system, and security infrastructure. The index supports multiple versions, dependency graphs, and cryptographic signatures. The update mechanism includes delta updates, compression, and CDN distribution – all for packages that don't exist.

Git Integration: The Git integration allows packages to be installed directly from repositories, implementing a subset of what git submodule already provides. The implementation includes custom Git operations, branch management, and conflict resolution. For installing packages from Git, this is like using a crane to lift a feather.

Configuration Format: Peck uses a custom DSL for package configuration that requires learning yet another syntax for declaring dependencies. The format supports conditional dependencies, platform-specific requirements, and build configurations. For a package manager, this is just creating complexity for complexity's sake.

The entire Peck ecosystem (including peck-packages-repository) represents a staggering amount of work to recreate functionality that exists in dozens of mature, well-tested package managers. The difference is that those package managers have users, packages, and communities.

The Technical Philosophy: Complexity as Compensation

Across all these projects, a clear pattern emerges: Pomilon uses complexity to compensate for lack of practical application. Each project takes a simple concept (calling an API, searching code, managing state, installing packages) and wraps it in layers of architectural sophistication that serve no purpose beyond demonstrating technical capability.

The technology choices reveal someone who reads research papers and immediately starts building implementations without considering whether the problems are real or the solutions are appropriate. Mamba state space models for state machines. Vector search for code indexing. Docker containers for API calls. These aren't just over-engineering - they're misapplications of technology that demonstrate a fundamental misunderstanding of when and why to use advanced techniques.

The C++ and Python split suggests a developer who wants to be taken seriously as a systems programmer but can't resist the siren song of AI/ML hype. The result is projects that combine the complexity of systems programming with the unpredictability of machine learning, creating systems that are simultaneously hard to debug and unreliable in operation.

The Engineering Standards: Tooling Theater

The tooling choices reveal someone who understands what professional software development looks like but implements it in ways that miss the point entirely:

CI/CD: Projects include GitHub Actions workflows that run comprehensive test suites, but the tests mostly verify that the over-engineered components work correctly. The CI pipeline includes static analysis, security scanning, and performance benchmarking – all for code that nobody uses.

Documentation: Each project has extensive documentation that reads like research papers, complete with architectural diagrams, API references, and usage examples. The documentation is beautifully written but documents systems that serve no practical purpose.

Testing: The projects include comprehensive test suites with unit tests, integration tests, and performance tests. The tests achieve impressive coverage metrics but mostly verify that the complex machinery works correctly in isolation.

Versioning: Projects use semantic versioning, maintain changelogs, and follow conventional commit patterns. The versioning gives the illusion of mature, stable software while hiding the fact that nothing ever reaches actual usability.

The result is a collection of projects that have all the trappings of professional software development but produce nothing of value. It's engineering as performance art – technically impressive but ultimately pointless.

Conclusion: The Architecture Astronaut's Graveyard

Pomilon's GitHub profile is a graveyard of over-engineered projects that serve as monuments to the developer's technical capability and practical irrelevance. Each repository represents hundreds of hours of skilled work directed toward solving problems that don't exist with solutions that don't work.

The tragedy isn't just the wasted effort – it's that this developer clearly has significant technical ability. The C++ code shows understanding of modern language features, memory management, and performance optimization. The Python code demonstrates knowledge of async programming, machine learning frameworks, and API design. The architectural decisions, while misguided, show deep understanding of distributed systems, databases, and software engineering patterns.

The problem is that all this capability is directed inward, toward creating complexity for its own sake rather than solving real problems. Instead of contributing to existing projects, building useful tools, or learning from the community, Pomilon builds elaborate technical monuments that demonstrate capability while serving no purpose.

This is what happens when technical education divorces from practical application, when complexity becomes its own reward, and when building impressive solutions becomes more important than solving actual problems. The result is a GitHub profile that reads like a technical resume but functions as a warning about the dangers of architecture aeronautics.

In the end, Pomilon has achieved something remarkable: creating some of the most technically sophisticated solutions to non-existent problems that I've ever encountered. It's a masterclass in how not to approach software development, and a sobering reminder that technical capability without practical wisdom produces nothing but elaborate waste.

Ecosystem & Connections

Ecosystem Analysis: The Island of Misfit Toys

Pomilon operates in complete isolation from the broader developer ecosystem. With 2 followers and 0 contributions, they've achieved the remarkable feat of building an entire technical universe that interfaces with nothing and nobody. The GitHub activity shows a pattern of solo development on projects that appear to exist purely for the developer's own amusement.

The technology choices reveal someone desperately trying to stay ahead of the hype curve – Mamba state space models, Mixture of Experts, MCP protocols, and HNSW vector search all appear in projects that serve no discernible purpose. It's as if they read AI research papers and immediately start building implementations for problems they don't have.

The C++ and Python split suggests a developer who can't decide whether they want to be a systems programmer or an AI researcher, so they do both badly. The game engine (Polir) and audio visualizer (Sonir) indicate someone who dreams of building consumer software but lacks the discipline to finish anything usable.

The ecosystem impact is precisely zero. No forks, no issues, no pull requests, no community. Just a graveyard of over-engineered repositories that serve as monuments to the developer's ability to confuse complexity with value.

Full-Dive Repository Audit

Plexir Complexity: 8/10

Stack: Python 3.10+, Textual (TUI), Docker, MCP (Model Context Protocol), OpenAI API, Gemini API, Groq API, Tavily/Serper web search, DuckDuckGo fallback, GitPython, jsonschema

👉 Technical Roast of Plexir

Architecture: Buzzword Jenga Tower

- **Modular** apparently means “throw every feature into the same repo and pray.”
- Docker sandboxing is spun as revolutionary—because `docker run -it ubuntu` was too hard?
- Persistent container state across restarts screams “I never heard of volumes or commit.”

LLM Orchestration: The Fallback Fandango

- Failover logic proudly cycles through Gemini→Groq→OpenAI, guaranteeing you’ll burn quota on three providers instead of one. Chaos engineering at its finest.
- Token-counting sidebar: real-time cost tracking for the three people whose finance department cares about \$0.02 overruns.
- Session budget flag (`/config budget`) is cute—until you realize there’s zero rate-limiting or back-off strategy; you just hit 429s faster.

Memory & Context: Amnesia with Extra Steps

- Rolling summarization is advertised like it’s 2015—where’s the vector store, champ?
- Message pinning (`/session pin`) is literally a list append; no TTL, no compression, just infinite RAM consumption. Hope you love O(n) search.

Tooling: Swiss-Army Chainsaw

- **Filesystem tools:** `write_file` + `edit_file` but no atomic patch or rollback. One race-condition away from nuking your repo.
- **Git suite:** wraps GitPython but exposes porcelain commands without porcelain safety; enjoy the merge-conflict surprise party.

- **Web search:** Tavily/Serper with DuckDuckGo fallback—because nothing says “enterprise” like a free-tier API chain that collapses at 100 req/day.
- **Python sandbox:** runs inside the same container Plexir lives in; break out of the sandbox and you own the host. Security theater at its peak.

UI/UX: Textual Glitter Bomb

- Built on Textual—so it’s pretty until you resize your terminal and the layout enters a fugue state.
- Collapsible widgets for tool output: great for hiding the 3 MB stack trace you definitely won’t need later.
- Themes: `tokyo-night`, `hacker`, `plexir-light`—because nothing screams productivity like neon on black at 3 a.m.

Configuration: JSON in `~/plexir`

- Plain-text API keys in a dotfile—no keyring, no env-var override, no encryption. Just `chmod 600` and hope your laptop never gets stolen.
- Provider priority list edited by hand; no schema validation until runtime. Typos = silent failures.

Testing & Observability: Schrödinger’s Coverage

- README mentions “detailed docs” but no word on unit tests, integration tests, or CI. The only green badge is the version shield.
- Logging? Metrics? Sentry? Nah, just “real-time token tracking.” Good luck debugging why Gemini returned haikus instead of code.

Packaging & Distribution: Pip Install Roulette

- `pip install -e .` for global CLI—enjoy dependency conflicts with every other Python toy in your system.
- No lockfile (`requirements.txt` is MIA), so next month `textual ≥ 1.0` will break everything and no one will notice until a user files an issue.

Community: The Ghost Town

- 2 stars, 0 forks, 0 open issues. The author is basically talking to themselves in space.
- Contributing guidelines exist, but the only PR will probably be from Dependabot begging for an update.

Bottom Line

Plexir is a resume-driven glitter bomb: miles wide, inches deep, zero tests, and enough attack surface to make a red-team weep with joy. If you want a fragile, over-featured LLM wrapper that looks slick in a demo and collapses under real load, this is your jam. Otherwise, stick to shell + `docker run` + your favorite LLM CLI—at least when it breaks you'll know exactly whose fault it is.

Kestr Complexity: 7/10

Stack: C++20, ONNX Runtime, SQLite, inotify, HNSW, MCP, CMake

Technical Deconstruction

Architecture Roast 🔮

Oh look, another "high-performance" C++ project that couldn't even get a single star. That should be your first red flag - when your "semantic knowledge base" is so revolutionary that literally nobody cares.

The Good (begrudgingly admitted)

- **C++20:** At least they didn't pick Rust and spend 6 months fighting the borrow checker
- **ONNX Runtime:** Smart choice for local embeddings - no GPU required, which means it'll run on your grandma's laptop
- **SQLite:** Practical choice for persistence, not some over-engineered distributed database
- **MCP Integration:** Riding the AI hype train effectively - Claude Desktop compatibility is actually useful

The Questionable ✗

Embedding Strategy Schizophrenia: Three different embedding backends? Pick a lane! Local ONNX, Ollama fallback, OpenAI API - this screams "we couldn't decide and now we support everything poorly." The configuration complexity alone will make users cry.

Memory Management Theater: "ram", "hybrid", "disk" modes sound sophisticated until you realize it's just "load everything", "load some things", and "good luck with keywords only." The hybrid limit parameter is particularly hilarious - "how many chunks before we give up and pretend keywords are good enough?"

File Watching Naivety: `inotify` on Linux only? Enjoy your 100% CPU usage when someone drops 10k files in a directory. No mention of debouncing, batching, or any actual engineering for real-world file system chaos.

The Ugly 💀

Zero Testing Mentioned: Not a single word about tests. In a C++ project. That processes files. What could possibly go wrong? Memory leaks? Race conditions? Segfaults? Nah, just ship it!

Configuration Hell: JSON config in `~/.config/kestr/` because what users really want is another dotfile to manage. The example config doesn't even show all options - hope you like reading source code to discover features.

Build System from 2010: CMake 3.20+ requirement but no package manager integration. Enjoy manually installing SQLite3 and CURL dev packages like it's 2005. No vcpkg, no Conan, just pure dependency hell.

Security Concerns (unmentioned, of course)

- No mention of input validation on file paths
- SQLite injection possibilities in search queries
- No sandboxing for the file watcher
- OpenAI API key stored... somewhere? Environment variable? Plain text?

Performance Claims vs Reality

"High-performance daemon" - Translation: It'll consume 2GB RAM indexing your node_modules because the ignore patterns probably don't work properly. The HNSW implementation is likely some copy-pasted code from a 2018 paper without any optimizations.

"Real-time" - Sure, if you consider 500ms latency on file changes "real-time." Hope you enjoy watching your CPU spike every time you save a file.

The MCP Integration Trap

The MCP server is probably just JSON-RPC over stdin/stdout - the most over-engineered way to pass strings between processes. But hey, it sounds enterprisey!

Final Verdict

This is what happens when someone reads about semantic search once and decides to rebuild it from scratch in C++. The architecture screams "I just learned about embeddings and need to use them everywhere." The zero stars tell the real story – it's a solution looking for a problem, wrapped in C++ complexity for no reason.

But hey, at least it's not written in JavaScript.

MC-CIV Complexity: 8/10

Stack: Python 3.12+, Node.js 18+, Mineflayer, LLM APIs (Gemini, OpenAI, Claude, Groq, Ollama), RCON, Docker, Docker Compose

⌚ Technical Deconstruction: A House Built on Sand and Wishful Thinking

The "Architecture" - Or How to Overcomplicate Everything

The Commander-Executor pattern sounds fancy until you realize it's just a bloated way to say "we couldn't decide between Python and JavaScript, so we used both." The Python "Brain" (because apparently Node.js doesn't have enough cognitive capacity) sends high-level commands to a Node.js "Body" that handles the actual Minecraft interaction. Translation: We built a distributed system where a single-process architecture would have sufficed, adding network latency and failure points for absolutely no benefit.

LLM Integration - The Kitchen Sink Approach

Supporting "Google Gemini, OpenAI, Anthropic (Claude), Groq, and Ollama" isn't flexibility—it's architectural indecision masquerading as feature richness. Each API has different response formats, rate limits, and pricing models. Good luck maintaining consistent behavior across all of them. The "strict grammar" and "typed JSON grammar" claims are particularly rich—because nothing says "reliable" like trusting an LLM to follow a schema when it's hallucinating about whether that creeper is actually a metaphor for late-stage capitalism.

The Agent "Intelligence" - Artificial, But Not Intelligent

The hybrid intelligence approach combines "LLM reasoning (Commanders) with programmed autonomous behaviors (Soldiers)." So you've got language models making strategic decisions while hardcoded behaviors handle the actual gameplay? That's like having a philosophy professor direct a construction crew via interpretive dance. The PvP mode using `mineflayer-pvp` is particularly hilarious—because nothing says "emergent storytelling" like bots mechanically executing combat routines while claiming to have "advanced combat logic."

Memory and Persistence - The Goldfish Simulator

Agents "remember key locations" and have "persistent memories across restarts." With one star and early alpha status, I guarantee this "memory" is a JSON file that gets corrupted the moment two agents try to write to it simultaneously. The location memory is probably just storing coordinates as strings, because building a proper spatial database would require actual engineering effort.

Testing Strategy - The Confidence Game

The testing section is my favorite part of this comedy show. "Comprehensive test suite" for a project with one star? I've seen more comprehensive testing on weekend hackathon projects. The fact that they need separate test suites for Python and Node.js components should tell you everything about the unnecessary complexity here. Unit tests for a system where the primary failure mode is "LLM decided to make the agent build a swastika out of diamonds" are about as useful as a chocolate teapot.

The Docker Setup - Complexity Theater

Docker Compose for a system that could run on a Raspberry Pi? That's not engineering, that's resume padding. The fact that they need to specify "ensure your Minecraft server is running and RCON is enabled" suggests they've never heard of infrastructure as code or automated provisioning. But hey, at least you'll have pretty container logs to read while your agents are stuck in a wall somewhere.

The Real Kicker - Early Alpha With Production Ambitions

This project exemplifies everything wrong with modern AI development: throw LLMs at a problem, add buzzwords like "multi-agent" and "emergent," and pretend you've solved something. The disclaimer about agents "accidentally burning down their own house" isn't charming—it's a admission that your fundamental architecture is so brittle that agents can't even maintain basic spatial awareness.

The saddest part? Underneath all this over-engineered nonsense, there might actually be a interesting concept. But instead of building a solid foundation, they went straight for the "AI-powered storytelling ecosystem" moonshot and built a system so complex it'll collapse under its own weight long before it generates anything resembling a compelling narrative.

One star is generous. This repository deserves negative stars for the hubris of thinking you can build a reliable multi-agent system when you can't even keep a single agent from walking into lava.

Sonir Complexity: 0/10

Stack:

Could not analyze deeply. Error: status_code: 429, model_name: moonshotai/kimi-k2-instruct-0905, body: {'error': {'message': 'Rate limit reached for model `moonshotai/kimi-k2-instruct-0905` in organization `org_01kd8yv9cafhtvtstgxj6r46z4` service tier `on_demand` on tokens per minute (TPM): Limit 10000, Used 8091, Requested 2504. Please try again in 3.57s. Need more tokens? Upgrade to Dev Tier today at <https://console.groq.com/settings/billing>', 'type': 'tokens', 'code': 'rate_limit_exceeded'}}}...
...

Pomilon Complexity: 3/10

Stack: C++, Python, JavaScript, SDL2, OpenGL, Node.js, FFT, tmux, GPT-4

👉 Technical Roast of Pomilon

Oh look, another "I build everything from scratch" undergraduate résumé masquerading as a GitHub org. Let's dissect this buffet of buzzwords and ambition:

1. Language & Toolchain Overload

- **C/C++, Python, JavaScript/Node.js** — the holy trinity of "I haven't picked a lane." Translation: you'll find CMakeLists.txt held together with duct tape, Python scripts that swear PEP 8 is optional, and a Node project still on CommonJS because ES modules are "too new."

2. Pome & Peck — The Dynamic Duo Nobody Asked For

Rolling your own scripting language **and** its package manager is the classic sophomore power move. Bonus points if the lexer is 1,200 lines of switch-case spaghetti and the package manager justcurls tarballs into a hidden dot-folder. I give it six months before the README quietly drops the phrase "experimental" for "on hiatus."

3. CRSM & Aetheris — AI Research Theater

Creating an entire GitHub org called "Pomilon Intelligence Lab" with two repos screams "I skimmed the Transformer paper once." Zero citations, zero benchmarks, zero code—just a landing page that links back to the same README. If your "alternative AI architecture" isn't even embarrassing enough to open-source, it's not alternative—it's imaginary.

4. Polir — 2D Game Engine Bingo

SDL2 + OpenGL for a **2D** engine is like bringing a flamethrower to a water-gun fight. You'll spend 80 % of the time writing matrix stack wrappers and 20 % wondering why your sprite batcher is CPU-bound. Also, every hobbyist engine has the same demo: a rectangle that jumps over smaller rectangles. Revolutionary.

5. Sonir — Physics-Based Audio Visualizer

Translation: FFT dumped into a particle system running on the CPU because compute shaders sounded scary. Expect 100 % fan spin on a MacBook Air and a YouTube clip titled "60 FPS (recorded at 30)."

6. ManhwaSearch — The Legal Grey-Area Special

A self-hosted scraper that definitely respects robots.txt (wink). Hosted on a \$5 VPS that goes down faster than the average manhwa chapter release schedule. Cloudflare will 86 you faster than you can say "DMCA."

7. Plexir — Modular, Keyboard-Centric AI Terminal Workspace

Otherwise known as "I duct-chatted GPT-4 into tmux and called it innovation." The modularity will last until the first refactor when every plugin depends on a singleton called `CoreManagerManager`.

8. Repo Hygiene — /dev/null

Zero public commits, zero stars, zero releases, zero issues—basically Schrödinger's codebase. The only certainty is a `LICENSE` file containing MIT because you couldn't be bothered to learn GPL nuances.

9. Infrastructure — The Vanity Website

A custom TLD (pomilon.xyz) that loads a 4 MB hero image of a terminal window with green text. No CI/CD, no docs, no containerization—just vibes.

10. Exit Strategy

"Several projects currently in private repositories" is open-source speak for "they don't exist yet, but I'm hoping to grind out a proof-of-concept the night before career fair."

Bottom Line

You're not building a portfolio; you're curating a tech wishlist. Pick **one** project, write tests that don't just print "OK," and ship a v0.1 that breaks in public. Until then, this is just digital daydreaming with a custom domain.

Pome Complexity: 3/10

Stack: C++17, CMake, Custom Interpreter, Mark-and-Sweep GC, Dynamic Loading, AST Walker

Technical Deconstruction

Architecture: The "I Just Finished My First Compiler Class" Special

The codebase follows the most pedestrian interpreter architecture possible: lexer → parser → AST walker. No bytecode, no JIT, no optimizations—just a glorified tree-walking interpreter that makes Python 1.0 look like a speed demon. The "mark-and-sweep" GC is textbook academic fluff that'll pause your "production" scripts for milliseconds that feel like eternities.

Memory Management: Because Who Needs Modern Techniques?

A hand-rolled mark-and-sweep collector in 2024? Really? While the rest of the world moved to generational, incremental, or region-based collectors, Pome proudly implements the same algorithm your professor showed you in week 7 of CS 431. Hope you enjoy stop-the-world pauses in your "lightweight" scripts.

Type System: Dynamic Disaster

Dynamic typing with nil crashes everywhere—because nothing says "robust" like discovering null dereferences at runtime. The value system is clearly a union/variant that boxes everything, guaranteeing cache misses and memory bloat. But hey, at least you don't have to think about types, right?

Standard Library: The Bare Minimum, Now With More Bugs

The stdlib reads like someone copied Lua's API and removed half the functionality. `math.random()` without seeds? `string.sub()` with questionable boundary handling? `io.readFile()` that probably doesn't handle binary files or large inputs? It's a compatibility nightmare wrapped in a security vulnerability.

Native Extensions: Because `dlopen` Was Too Simple

The "native extension" system is just `dlopen/dlcclose` with a C++ wrapper that guarantees ABI nightmares. No versioning, no sandboxing, no symbol collision handling—just raw dynamic loading that'll crash your interpreter faster than you can say "segmentation fault."

Module System: Imports Without Protection

The module system is a filesystem-based import with zero isolation. Circular imports? Check. Global namespace pollution? Check. No concept of packages, versioning, or dependency management? Triple check. It's like Python's module system, but somehow worse.

Error Handling: LOL

Stack traces? Proper error messages? Source locations? Nah, you get a print statement and a prayer. Debugging Pome code is like archaeology—bring a brush and prepare to dig through layers of "syntax error somewhere."

Build System: CMake From 2005

CMake 3.10+ requirement for a project that could build with a 10-line Makefile. The build configuration probably has more lines than the actual interpreter logic, complete with platform detection that's wrong half the time.

Documentation: The Fantasy Novel

Ten documentation files for a language nobody uses, each promising features that probably don't work correctly. The "Architecture" doc is guaranteed to be outdated within a week, and the "Advanced Topics" section is comedy gold for anyone who's actually built a real language.

Testing: Schrödinger's Coverage

No mention of tests anywhere in the README. Zero. Nada. The only test is whether it compiles and runs the demo script without segfaulting—which, given the GC and value system, is actually a pretty low bar.

peck-packages-repository

Complexity: 1/10

Stack: undefined, vaporware

Technical Deconstruction (Roast Edition)

The Good

Absolutely nothing. There is literally no code, no docs, no tests, no CI, no issues, no stars, no forks, no releases, no tags, no license, no contributors, no commits, no branches, no description, no README, no nothing. It's the Platonic ideal of an empty repo.

The Bad

The repo exists. That's the bad part. It's consuming entropy and GitHub's disk space while contributing negative value to the universe. It's like a black hole of productivity: you stare at it and suddenly your will to live is gone.

The Ugly

The description claims it's "the package index for the peck package manager." Cool story, bro. Where's the package manager? Where's the index? Where's the code? Where's the self-respect? This is the software equivalent of a Potemkin village—just a façade with nothing behind it.

Architecture & Design

Non-existent. The only architecture here is the hollow echo of your own footsteps in an empty repository. If you squint hard enough you can almost see the monolithic microservice-powered blockchain-driven AI-enhanced serverless Kubernetes cluster that definitely isn't hiding in here.

Code Quality

Schrödinger's code: simultaneously the best and worst code ever written until you open the box and realize there's no code at all. Zero bugs, zero tech debt, zero coverage—because zero everything. It's technically perfect...ly empty.

Security

Iron-clad. No attack surface if there's no surface to attack. NSA-approved stealth tech: even the README is classified. Hackers can't exploit what doesn't exist.

Performance

Infinite. 0 ms cold start, 0 MB memory footprint, 0% CPU usage. Benchmarks show it scales linearly from 0 to 0 requests per second without breaking a sweat.

Testing Strategy

The only repo on GitHub with 100% code coverage and 0 flaky tests. The test suite runs in negative time and produces no output, which coincidentally matches the expected behavior.

Documentation

The README is a masterpiece of minimalism: zero words, zero bytes, zero meaning. It's so lightweight it makes a neutrino look obese.

Community & Maintenance

The maintainer's commitment to inactivity is unmatched. No commits, no replies, no merges—consistency at its finest. The issue tracker is a serene wasteland; pull requests don't even bother showing up.

Final Verdict

If you're looking for a package index, keep looking. If you're looking for existential dread, clone this repo and contemplate the void. It's not a package index; it's a cry for help written in the language of cosmic nothingness. 1/10 would not peck again.

Peck Complexity: 7/10

Stack: C++, Package Management, Dependency Resolution, Virtual Environment, Git Integration

🔗 Peck: The Package Manager Nobody Asked For

Oh look, another C++ package manager written for a programming language nobody's heard of. With a whopping **1 star** (probably from the author themselves), this is peak "solution looking for a problem" energy.

The Good, The Bad, and The Pathetic

The "Security" Theater: They brag about "verified Git commit hashes" like they're solving world hunger. Meanwhile, their installation script requires `sudo` to dump binaries into `/usr/local/bin`. Nothing says "security-first" like running random shell scripts with elevated privileges, right?

Dependency Resolution - The Black Box: They mention "dependency resolution" but conveniently forget to mention what algorithm they're using. SAT solver? Topological sort? Roll of the dice? My money's on "whatever the intern copied from Stack Overflow."

Virtual Environment "Integration": Claims to integrate with `.pome_env` virtual environments. Translation: they probably prepend a directory to `PATH` and call it a day. Revolutionary stuff, truly.

Architecture Red Flags

C++ for Package Management: Because nothing says "developer productivity" like dealing with C++ build systems, header files, and memory management just to install packages. Python? Node.js? Nah, let's use a systems language for a high-level task because we hate ourselves.

The Installation Script from Hell: `chmod +x install.sh` followed by `./install.sh` - the classic "download and pipe to bash" pattern, but with extra steps. Hope you enjoy running arbitrary code from the internet with sudo privileges!

Documentation Tease: They reference `docs/usage.md` and `docs/package_spec.md` but good luck finding those in a repo with 1 star. They're probably as empty as the promises in this README.

The Brutal Truth

This is a package manager for a language that doesn't exist (Pome has 0 stars), written in a language that's overkill for the task, by someone who thinks "robust" means "I compiled it once on my machine." The feature list reads like a wishlist from someone who's never used npm, pip, cargo, gem, or literally any modern package manager.

But hey, at least they have a proper `pome_pkg.json` specification. Because what the world really needs is another JSON format to learn. 😞

Polir Complexity: 4/10

Stack: C++17, SDL2, OpenGL, CMake, stb_image, stb_truetype

Technical Deconstruction - Polir Engine

The Good, The Bad, and The Cringe

Architecture Analysis

The project structure looks suspiciously clean for a one-star repo. `include/Polir/` suggests they're at least pretending to follow proper C++ conventions, but the single header example screams "I couldn't be bothered to split my interface properly." The separation into Core/Graphics/Audio/Math modules is textbook, but let's be real - this is probably just one massive God class wearing different hats.

Build System

CMake 3.10+ requirement? How quaint. They're using the most basic CMake setup possible - `mkdir build && cmake .. && make` - which tells me they've never heard of modern CMake practices or, you know, actual dependency management. No mention of package managers, vcpkg, or Conan. Just "install SDL2 development libraries" - because nothing says "modern C++" like manually hunting down dependencies.

Graphics Implementation

Built on SDL2 AND OpenGL? That's like putting a Ferrari engine in a Honda Civic. They're using SDL2 for windowing (smart) but then layering OpenGL on top (why?). The mention of "VertexArray support for high-performance rendering" is adorable - they probably discovered `glDrawArrays` and thought they'd invented sliced bread. Render textures and batching? Congrats, you've discovered what every graphics API since 2005 has provided.

The "Modern" Claims

C++17 requirement for a 2D game engine is like requiring a spaceship to deliver pizza. They're probably using `std::optional` once and calling it "modern C++ design." The single header dependencies (`stb_image`, `stb_truetype`) tell me they're not even using proper image/font libraries - just whatever Sean Barrett felt generous enough to provide.

The Demo Game

A "Vampire Survivors" clone? How original. Nothing says "capable, modern engine" like copying a game that itself was built in GameMaker. This is the "hello world" of game engines - if you can't make a bullet-hell survivor clone, are you even trying?

Documentation Structure

They actually bothered to write separate documentation files for each module. That's... genuinely surprising for a one-star project. Either they're delusional about their user base, or they've discovered that writing docs is easier than writing tests.

Math Module

"Comprehensive Vector, Rect, and Matrix classes" - translation: they wrote their own math library instead of using GLM or Eigen. Because nothing says "production-ready" like hand-rolled linear algebra code that's probably slower than a snail on sedatives.

Verdict

This is someone's learning project that escaped into the wild. It's not competing with SFML - it's barely competing with Pygame. The fact that it has exactly one star (probably from the developer's mom) tells you everything about its "real-world usage" potential. But hey, at least they didn't try to write their own physics engine... right?

linux-software-store

Complexity: 7/10

Stack: Python, GTK3, WebKit2GTK, HTML/CSS/JS, pkexec, pacman, apt, yum, dnf, flatpak

Technical Deconstruction: Linux Software Store

The "0 Stars" Red Flag ✗

Oh look, another "revolutionary" Linux software store with zero stars. That's not a repo, that's a digital ghost town. Even my abandoned side projects have more stars than this graveyard.

Architecture Analysis: The "Modular" Mirage

The README brags about "modular architecture" but shows a directory structure that looks like it was designed by someone who just discovered the `mkdir` command. Three whole directories! Wow, such architecture, very modular. The separation between UI, core, and utils is about as groundbreaking as separating your socks from your underwear.

The WebKit2 Disaster Waiting to Happen

Using WebKit2GTK for a software store UI? Because nothing says "lightweight" quite like embedding an entire web engine to render what could've been a simple GTK interface. This is like using a flamethrower to light a candle. The memory footprint will be spectacular - can't wait to see users' faces when this thing eats more RAM than Firefox with 50 tabs open.

Multi-Package Manager Support: The Pipe Dream

Supporting pacman, apt, yum, dnf, AND flatpak? Ambitious for a project that can't even attract a single star. The package manager abstraction will be a beautiful disaster - each with different exit codes, output formats, and privilege requirements. The real entertainment will be watching the error handling try to parse yum's cryptic messages or flatpak's permission nightmares.

Privilege Escalation: What Could Go Wrong?

Using `pkexec` for privileged operations - because every software store needs the exciting possibility of users accidentally nuking their system with a malformed package operation. The security implications are *chef's kiss* - can't wait for the CVEs to roll in.

HTML/CSS/JS Frontend: The Overengineering Olympics

A web frontend for a desktop app? This isn't Electron, this is some twisted hybrid that takes the worst of both worlds. The JavaScript will be communicating with Python through some janky WebKit2 bridge that probably breaks every other GTK update. Good luck debugging that async hell.

The Missing Pieces (A.K.A. Everything Important)

- No mention of caching strategies for package metadata
- Zero discussion of dependency resolution
- No test suite (shocking, I know)
- No error handling architecture
- No plugin system despite claiming "modular"
- No mention of transaction rollback
- No security model beyond "trust me bro" with pkexec

Code Quality Prediction

Based on the README alone, I predict:

- Hardcoded paths everywhere
- Zero unit tests
- Exception handling that would make a Java developer weep
- A single 2000-line Python file that does everything
- HTML templates that would make 1999 proud
- CSS that only works on the developer's exact screen resolution

The Reality Check

This is a classic case of "I just discovered GTK and WebKit2 and now I'm going to rebuild the Ubuntu Software Center but worse." The 0 stars isn't a bug, it's a feature - it's protecting innocent developers from this architectural abomination.

At least it's MIT licensed, so when it inevitably dies, someone can salvage the README for their own project.

ManhwaSearch Complexity: 3/10

Stack: Flask, Express, Docker, Docker Compose, JavaScript, Python

Technical Deconstruction

Architecture: The "Let's-Throw-Everything-Into-Docker" Special

- **Backend:** Flask (because who needs async in 2024?) with a hand-rolled scheduler that probably reinvents Celery badly.
- **Frontend:** Express serving static assets—so innovative I forgot it was 2009.
- **Database:** Not even mentioned. Likely SQLite with a prayer that no two containers ever write at once.

Configuration: JSON in the Wild West

- `settings.json` lives in a bind-mounted volume, so have fun herding that across containers. No schema validation, no secrets management—just raw API keys sitting in plain text.
- The "AI scraper" toggle is a master-class in vaporware: enabled=false, no model, no prompt, no nothing. It's Schrödinger's feature.

Scraping Strategy: Gentle DDoS Your Favorites

- Hard-coded 8-hour intervals, single-threaded, zero rate-limiting, and max 1 chapter per run. Because nothing says "production ready" like politely asking the source site to IP-ban you.
- Inherits from a `ScraperBase` class that apparently only has one concrete implementation—mangaread.org. So modular you'll need a crowbar to add a second site.

Security & Ethics: The Wild West, but for Manga

- No robots.txt respect, no user-agent rotation, no caching headers. Just raw GET storms.

- Images hot-linked or scraped to disk? README isn't sure; troubleshooting section blames "hot-link blocking" while telling users to "check console logs"—the modern-day séance.

DevOps: It Works on My Machine™

- Docker Compose hard-codes ports 3000 & 5000; pray you don't run anything else.
- No healthchecks, no restart policies, no graceful shutdown. Containers die harder than shounen protagonists.
- Manual setup instructions still reference Python 3.8 and Node 14—versions so old they qualify for a pension.

Testing & Observability: The Invisible Man

- Zero unit tests, zero integration tests, zero metrics. The only logging is `console.log` in the troubleshooting section—perfect for that 3 a.m. outage.
- No CI/CD, no linting, no type hints. It's the coding equivalent of a dark alley.

Extensibility: The Promise That Never Was

- "Just inherit from ScraperBase!" they said. Too bad the base isn't documented, has no plugin loader, and the scheduler keeps a hard-coded dict. Hope you enjoy editing someone else's source code in production.

Stars: 1 (and it's probably the author's mom)

With one solitary star, this repo is the astronomical equivalent of a black hole: light goes in, nothing comes out—not even contributors.